



Experiment No. 7
Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:



CSL701: Machine Learning Lab

Aim: Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, perform dimensionality reduction on the given dataset and maximize the accuracy, Precision, Recall, F1 score. **Theory:**

In machine learning classification problems, there are often too many factors on the basis of which the final classification is done. These factors are basically variables called features. The higher the number of features, the harder it gets to visualize the training set and then work on it. Sometimes, most of these features are correlated, and hence redundant. This is where dimensionality reduction algorithms come into play. Dimensionality reduction is the process of reducing the number of random variables under consideration, by obtaining a set of principal variables. It can be divided into feature selection and feature extraction. **Dataset:**

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute

Information: Listing of

attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

fnlwgt: continuous.

CSL701: Machine Learning Lab

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black. sex: Female,

Male. capital-gain: continuous. capital-loss: continuous. hours-per-week:

continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

CSL701: Machine Learning Lab

```
In [2]: import numpy as np
import pandas as pd

df = pd.read_csv("adult.csv")
df.head()
```

```
Out[2]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female

```
In [3]: df.describe()
```

```
Out[3]:
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.000000
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.422382
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.391444
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.490400e+06	16.000000	99999.000000	4356.000000	99.000000

```
In [4]: df.shape
Out[4]: (48842, 15)
```

```
:
In [5]: df.info
```

```
Out[5]: <bound method DataFrame.info of
on educational-num \
0      25      Private  226802      11th      7
1      38      Private  89814      HS-grad      9
2      28      Local-gov  336951      Assoc-acdm      12
3      44      Private  160323      Some-college      10
4      18      ?      103497      Some-college      10
...      ...      ...      ...      ...      ...
48837  27      Private  257302      Assoc-acdm      12
48838  40      Private  154374      HS-grad      9
48839  58      Private  151910      HS-grad      9
48840  22      Private  201490      HS-grad      9
48841  52      Self-emp-inc  287927      HS-grad      9

      marital-status      occupation      relationship      race      gender \
0      Never-married      Machine-op-inspct      Own-child      Black      Male
1      Married-civ-spouse      Farming-fishing      Husband      White      Male
2      Married-civ-spouse      Protective-serv      Husband      White      Male
3      Married-civ-spouse      Machine-op-inspct      Husband      Black      Male
4      Never-married      ?      Own-child      White      Female
48837      Tech-support      wife      White      Female
48838      Married-civ-spouse      Machine-op-inspct      Husband      White      Male
48840      Married-civ-spouse      Adm-clerical      Husband      White      Male
48841      Never-married      Adm-clerical      Own-child      White      Female
48841      Married-civ-spouse      Exec-managerial      wife      White      Female

      capital-gain      capital-loss      hours-per-week      native-country      income
0      0      0      40      United-States      <=50K
1      0      0      40      United-States      <=50K
2      0      0      40      United-States      <=50K
3      7688      0      40      United-States      >50K
4      0      0      30      United-States      <=50K
...      ...      ...      ...      ...      ...
48837      0      0      38      United-States      <=50K
48838      0      0      40      United-States      >50K
48839      0      0      40      United-States      <=50K
48840      0      0      20      United-States      <=50K
48841      15024      0      40      United-States      >50K

[48842 rows x 15 columns]>
```

```
:  
In [7]: df.isnull().sum()
```

```
Out[7]: age                0  
workclass            2799  
fnlwgt              0  
education            0  
educational-num      0  
marital-status       0  
  
occupation          2809  
relationship         0  
race                 0  
gender              0  
capital-gain         0  
capital-loss         0  
hours-per-week       0  
native-country       857  
income              0  
dtype: int64
```

```
In [9]: for col in ['workclass', 'occupation', 'native-country']:  
        df[col].fillna(df[col].mode()[0], inplace=True)  
df.isnull().sum()
```

```
Out[9]: age                0  
workclass            0  
fnlwgt              0  
education            0  
educational-num      0  
marital-status       0  
  
occupation          0  
relationship         0  
race                 0  
gender              0  
capital-gain         0  
capital-loss         0  
hours-per-week       0  
native-country       0  
income              0  
dtype: int64
```

```
:  
In [11]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, ran
```

```
:  
In [14]: from sklearn import preprocessing  
categorical = ['workclass', 'education', 'marital-status', 'occupation', 'rela  
for feature in categorical:  
    label = preprocessing.LabelEncoder()  
    X_train[feature] = label.fit_transform(X_train[feature])  
    X_test[feature] = label.transform(X_test[feature])
```

```

In [15]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
X_train.head()

```

```

Out[15]:

```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship
0	-0.849978	-1.887643	-0.551219	1.212393	-0.027733	-0.406325	-1.554732	0.969833
1	0.241031	-0.094859	1.687545	-2.650223	-1.587187	-0.406325	-1.049322	0.969833
2	-0.486308	1.697924	-1.434052	-0.590161	0.362131	-0.406325	-0.543912	-0.899325
3	-0.195373	-0.094859	-0.384485	1.212393	-0.027733	0.922720	-0.796617	-0.276272
4	-0.704510	-0.094859	1.608144	0.182362	-0.417596	1.587242	1.730434	1.592886

```

In [18]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

LR = LogisticRegression()
LR.fit(X_train, y_train)

```

```

Out[18]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or
trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page
with nbviewer.org.

```

```

In [20]: y_pred = LR.predict(X_test)
accuracy_score(y_test, y_pred)

```

```

Out[20]: 0.8221524602470484

```

```

In [21]: from sklearn.decomposition import PCA
pca = PCA()

```

```

In [22]: X_train = pca.fit_transform(X_train)
pca.explained_variance_ratio_

```

```

Out[22]: array([0.14740223, 0.10130193, 0.08096753, 0.07933632, 0.07433976,
0.07314763, 0.07066221, 0.06753572, 0.06516078, 0.06093536,
0.06003764, 0.04864317, 0.04289137, 0.02763835])

```

```

In [24]: X = df.drop(['income'], axis=1)
y = df['income']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, ran

```

```
:
In [25]: categorical = ['workclass', 'education', 'marital-status', 'occupation', 'rela
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])
```

```
In [26]: X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
```

```
In [27]: pca= PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
dim = np.argmax(cumsum >= 0.90) + 1
print('The number of dimensions required to preserve 90% of variance is',dim)
```

The number of dimensions required to preserve 90% of variance is 12

```
In [28]: X = df.drop(['income', 'native-country', 'hours-per-week'], axis=1)
y = df['income']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, ran
```

```
In [30]: categorical = ['workclass', 'education', 'marital-status', 'occupation', 'rela
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
```

```
In [31]: LR2 = LogisticRegression()
LR2.fit(X_train, y_train)
```

```
Out[31]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
:
In [32]: y_pred = LR2.predict(X_test)
accuracy_score(y_test, y_pred)
```

```
Out[32]: 0.8229031597625059
```



```

In [33]: from sklearn.metrics import confusion_matrix
import pandas as pd
confusion = confusion_matrix(y_test, y_pred)
df_confusion = pd.DataFrame(confusion, columns=['Predicted No', 'Predicted Yes'])
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
<=50K	0.84	0.94	0.89	11138
>50K	0.71	0.44	0.54	3515
accuracy			0.82	14653
macro avg	0.78	0.69	0.72	14653
weighted avg	0.81	0.82	0.81	14653



Conclusion:

1. Using principal component analysis, we were able to obtain an accuracy score of 0.82 on the testing data, meaning that our model is 82% accurate.
2. Our model's precision score, which quantifies the accuracy of the positive predictions, is 0.84.
3. Recall evaluates the model's capacity to accurately identify all pertinent instances; our model received a recall score of 0.94.
4. The F1-score, which our model yielded, is 0.89. It is the harmonic mean of precision and recall and offers a balance between the two metrics.